

URI SOCIAL

Image Generation

Master Rulebook V3

Revised After Expert Review

Sharper vocabulary. Hard exclusions. Deeper compositing.

May 2026 — URI Creative Limited

0. What Changed From V2

This V3 incorporates expert review feedback. Major changes:

Change	Reason
Expanded aesthetic vocabulary with missing cinematography, product photography, and fashion clusters	V2 was too dependent on a small set of overused tokens (cinematic, editorial, premium)
Deprecated the “aesthetic-first dominates” framing	Mechanistically inaccurate. Transformers attend globally.
Replaced with “semantic coherence + entropy reduction” framing	More accurate; informs why short, uninterrupted prompts work
Operationalized “immersive” with concrete vocabulary	Was too vague; the model has no stable internal definition of ‘immersive’
Added perspective matching, ambient occlusion, and material-aware relighting to compositing	V2 covered lighting and reflections but missed camera angle, focal length, vanishing points, and material-specific physics
Converted sensitive content protection from -200 penalty to HARD EXCLUSION	Penalties are weak; exclusion logic is what safety actually requires
Added ensemble moderation and manual escalation for sensitive content	Single-classifier dependency is fragile
Expanded style ontology from 3 dimensions (intent, emotion, energy) to 11 dimensions	V2 was shallow; visual branding has higher dimensionality
Deepened African realism rules with specific local references	‘Natural skin tones’ alone produces Western stock photography with dark skin
Cut diluted internet-aesthetic tokens	‘Award-winning,’ ‘masterpiece,’ ‘ultra-detailed,’ excessive ‘cinematic’ — all signal noise

1. Context for the Reviewer

This document consolidates every rule, principle, and constraint governing image generation in URI Social. It is intended for engineering implementation and external expert review.

What URI Social Is

URI Social is an AI-powered social media management product for African small businesses, primarily in Nigeria. The core AI agent is called Jane. End users (business owners) interact with Jane mainly through WhatsApp. They send Jane a request or upload a product photo, and Jane generates marketing images, captions, videos, and blog content across Instagram, Facebook, LinkedIn, X, TikTok, and WhatsApp Status.

What Powers Image Generation

- Primary image model: GPT Image 2 (gpt-image-1 in the API)
- Classification and vision analysis: GPT-4o-mini
- Conversational reasoning: Claude Sonnet 4
- Video generation: Seedance and Kling
- Compositing and final rendering: Sharp + Canvas (Node.js server-side)

The Commercial Constraints

- Client products (lipstick, food dish, packaged good) must be preserved exactly. No AI-redesign.
- Brand colours must appear naturally, not as a flat overlay.
- Text and logos must be controllable, with no gibberish.
- Output must be reliable across an entire campaign, not just good once.
- Visual style must match the user's brand voice AND the content's emotional context.

2. The Core Philosophy

2.1 The Fundamental Principle

THE PRINCIPLE

Activate the model's best training. Then protect the brand. High-coherence aesthetic priors should remain uninterrupted by low-aesthetic operational constraints.

2.2 The Updated Mental Model

Image models match patterns from training data more than they follow procedural instructions. But it is **INCORRECT** to say the model doesn't reason. Modern multimodal models do perform spatial, compositional, and constraint-satisfaction reasoning. They do it weakly compared to latent prior matching, but they do it.

The right framing is: procedural instructions fail when they describe IMPLEMENTATION DETAILS. They work when they describe PERCEPTUAL OUTCOMES.

Bad (implementation detail)	Good (perceptual outcome)
Use three distinct depth layers	Cinematic depth separation with soft foreground occlusion
Apply 45-degree light from upper left	Warm directional side-light with motivated shadow direction
Ensure tonal continuation across the frame	Rich tonal range, atmospheric haze, shallow focus
Position product at the rule of thirds intersection	Editorial composition with the product as visual anchor

2.3 The Dilution Principle (Mechanistically Accurate)

Long constraint blocks degrade output quality. This is **NOT** because earlier tokens dominate — transformers attend globally. The actual mechanism is:

- Semantic conflict between competing distributions
- Entropy from too many directives
- Coherence loss when literal constraints interrupt aesthetic priors

Therefore: keep constraints compact, minimal, and semantically distant from aesthetic blocks. Where possible, move constraints OUT of the prompt entirely into the compositing layer.

3. The Prompt Architecture

Every image prompt is assembled in this exact order. Aesthetic blocks come first, kept dense and coherent. Constraints come last, kept tight and semantically isolated.

Order	Block	Purpose	Length
1	Quality Benchmark	Puts the model in the right quality neighbourhood	1-2 lines
2	Genre & Style	The aesthetic genre and feel	2-3 lines
3	Composition	Framing and hierarchy in evocative terms	1-2 lines
4	Subject	Who or what is in the image, with realism cues	2-3 lines
5	Environment	The setting, lived-in and real	1-2 lines
6	Lighting	Mood and direction in cinematography vocabulary	1-2 lines
7	Texture & Color	Film grain, palette, grading	1-2 lines
8	Mood	Emotional register — 2-3 words max	1 line
9	[CONSTRAINT] Product	Product accuracy — stated tightly	1-2 lines
10	[CONSTRAINT] Brand & Text	Brand colours, text limits	1-2 lines

3.1 Full Assembled Example

Premium product campaign visual. Editorial commercial photography, shot on medium format. Magazine cover quality. Luxury beauty editorial. Sensual, atmospheric, with editorial retouching and luxury campaign photography standards. Centered editorial framing. Product as visual anchor with foreground occlusion and environmental wrapping. [PRODUCT – the provided lipstick] grounded naturally in the scene. Atmospheric studio environment with floating petals and soft mist. Warm tungsten practical lighting. Soft rim light, motivated shadow direction, soft halation. Luminous highlights with controlled specular reflections. 35mm film grain. Atmospheric haze. Kodak Portra tones. Rich saturated palette. Tactile material detail with photoreal material response. Mood: sensual, intimate. PRODUCT: Use the uploaded lipstick as the exact canonical reference. Preserve packaging geometry, label layout, typography, proportions, and material finish identically. BRAND: Subtly integrate deep magenta (#C2185B) through environmental accents, props, reflections, or lighting cues. TEXT: Leave clean negative space upper third. Render NO text in the image; typography will be composited in post-processing.

4. The Aesthetic Vocabulary

This is the high-signal vocabulary library, expanded with the missing clusters identified in expert review.

4.1 Quality Benchmarks (Block 1)

Naming a quality benchmark and a specific camera reference tells the model what tier of image to produce. Camera references are among the highest-signal phrases available.

- Premium editorial photography / editorial campaign / magazine cover quality
- Commercial product photography / advertising campaign
- Vogue-style / runway editorial / luxury campaign photography
- Documentary realism / film still / cinematic campaign
- Shot on 85mm / shot on medium format / shot on Hasselblad
- Shot on Phase One / shot on Mamiya 645 (medium format alternatives)

4.2 Cinematography Lighting Cluster (Block 6) — EXPANDED

Replace generic “cinematic light” with specific cinematography vocabulary:

- Tungsten practical lighting / practical tungsten glow
- Anamorphic lens flare / soft halation
- Motivated lighting / motivated shadow direction
- Soft diffusion bloom / black-pro mist diffusion / diffusion-filter softness
- Edge lighting / rim lighting / overexposed highlight bloom
- Negative fill / soft bounce lighting
- Low-key portrait lighting / chiaroscuro
- Volumetric light / atmospheric backlight
- Golden-hour atmosphere / warm directional side-light
- Overcast natural diffusion / soft window light
- Cinematic skin rendering / soft skin tones in available light
- Handheld documentary realism (for less polished, more authentic scenes)

4.3 Product Photography Cluster — NEW

Specifically for product-centric shots, these activate commercial product photography priors:

- Controlled specular reflections / specular highlight control
- Precision edge lighting / luxury packshot lighting
- Reflective surface management / controlled studio reflections
- Glossy acrylic reflections / glass refraction accuracy
- Liquid translucency realism / physically accurate condensation
- Macro texture realism / premium material rendering
- Photoreal material response (for plastic, metal, fabric, glass)
- Suspended liquid dynamics (for splash, pour, droplet effects)

4.4 Fashion / Editorial Cluster — NEW

- Off-camera flash realism / off-camera flash aesthetic

- Runway editorial energy / runway photography
- Street-style editorial / urban editorial framing
- Editorial retouching / high-end retouching
- Magazine spread framing / editorial layout
- Brutalist editorial layout / minimalist editorial
- Campaign photography / luxury beauty campaign

4.5 Texture & Atmosphere (Block 7) — EXPANDED

- 35mm film grain / 120 medium format grain / analog texture
- Kodak Portra look / Kodak Vision3 grading / Fuji 400H tones / Fuji cinematic grading
- Edge-rolloff softness / soft highlight bloom
- Natural environmental imperfections / lived-in realism
- Atmospheric haze / bokeh / shallow depth of field
- Cinematic color grading / film-stock grading
- Tactile surface detail / material physics

4.6 Color Systems (Block 7)

- Warm muted cinematic palette / earthy tones
- Rich saturated colour / vibrant editorial palette
- Monochromatic / desaturated luxe / tonal limited palette
- High contrast / deep shadows / luminous highlights
- Teal-orange grading (use sparingly — has become an aesthetic cliché)

4.7 Mood (Block 8 — single words, MAX 2-3)

Limit mood words to 2-3. Spamming mood words is dilution. One or two specific terms beat five generic ones.

- Premium contexts: elegant, refined, intimate
- Energetic contexts: vibrant, bold, playful
- Calm contexts: serene, grounded, natural
- Startup/professional: driven, authentic

5. Banned and Weak Vocabulary

Updated based on expert review. These words appear in many AI prompts but contribute little signal or actively dilute output quality.

5.1 Diluted Internet-Aesthetic Tokens (CUT)

So overused in AI prompts that they've lost signal value:

- “Award-winning”
- “Masterpiece”
- “Ultra-detailed” (use “commercial photographic realism” instead)
- “Hyperrealistic” (can produce plastic CGI artifacts — use “photoreal material response”)

5.2 Excessive Use of Common Tokens (LIMIT)

These are not banned, but limit their use. Better, more specific alternatives exist:

Overused Token	Better Alternative
cinematic	tungsten practicals, anamorphic, handheld documentary, soft halation
editorial	magazine spread, runway editorial, brutalist editorial, campaign photography
premium	luxury packshot, editorial retouching, controlled diffusion
dramatic	chiaroscuro, low-key, motivated shadow direction
beautiful	specific perceptual outcome (e.g., soft skin rendering, luminous highlights)

5.3 Weak Mood Words When Spammed (USE SPARINGLY)

Single use is fine. Spamming kills the signal. Cap at 2-3 mood words per prompt.

- Premium, elegant, refined, timeless — use ONE, not all four
- Driven, human, authentic, aspirational — same
- Aspirational — too broad on its own

5.4 Implementation-Detail Language (CUT)

Procedural instructions about pixels, layers, or technical execution — the model doesn't reason about these:

- “Three depth layers”
- “45-degree directional light”
- “Rule-of-thirds at exact coordinates”
- “Surface continuation”
- “Tonal continuation across the frame”

5.5 Vague Abstractions (CUT or OPERATIONALIZE)

- “Human ambition” — semantic abstraction, no visual mapping
- “Emotionally grounded” — weak visual mapping
- “Authentic energy” — fuzzy

- “Grounded naturally” — weak
- “Story-centered composition” — vague (replace with specific framing language)

6. Operationalizing “Immersive”

The term “immersive composition” was widely used in V2 but the model has no stable internal definition of it. V3 replaces it with concrete vocabulary.

6.1 What “Immersive” Actually Means

When we say a composition should be immersive, we mean specifically:

- Environmental wrapping — the scene continues around the subject in all directions, not just behind
- Foreground occlusion — elements in front of the subject partially obscure it, creating depth
- Atmospheric layering — visible distinction between near, mid, and far elements
- Environmental depth cues — atmospheric haze, light scattering, or environmental color cast
- Spatial continuity — the subject sits in the space, not on top of it

6.2 Vocabulary To Use Instead of “Immersive”

Instead of	Use
“immersive composition”	“environmental wrapping with foreground occlusion”
“immersed in the scene”	“grounded with atmospheric layering and spatial continuity”
“dreamlike immersion”	“atmospheric haze with environmental color bounce”
“360-degree environment”	“wrapping environment with environmental depth cues”

6.3 When To Use This Pattern

- Default for product hero shots in atmospheric styles
- All luxury and editorial campaigns
- Lifestyle and brand-story content
- NOT used for: clean studio packshots, minimal compositions, meme content, sale graphics with text zones

7. Product Preservation (Expanded)

This is the most critical commercial constraint. V3 expands the compositing pipeline with three additions identified by expert review: perspective matching, surface interaction simulation, and material-aware relighting.

7.1 The Two-Pass Compositing Pipeline

For product-centric brands (cosmetics, FMCG, packaged goods), the product should NOT be AI-generated. The pipeline:

1. Classify the product material (plastic, metallic, glass, fabric, ceramic, translucent) using GPT-4o-mini vision
2. Generate the SCENE with a product-shaped placeholder. Specify the desired camera angle and focal length explicitly
3. Segment the real product from the uploaded photo
4. PERSPECTIVE MATCH: detect the scene's camera angle and focal length, then warp the product cutout to match. Vanishing points must align.
5. MATERIAL-AWARE RELIGHT: apply lighting model appropriate to the product's material class
6. SURFACE INTERACTION: synthesize ambient occlusion at contact points, environmental bounce color on the product, and contact reflections
7. Composite the relit, warped product into the scene
8. Final low-strength image-to-image pass over the product-and-surroundings to marry the layers

7.2 Perspective Matching — NEW IN V3

The single biggest gap in V2 was that compositing focused on lighting and shadows but ignored geometric correctness.

- CAMERA ANGLE MATCHING — detect the scene's camera angle (low, eye-level, high, top-down). Warp the product cutout to match.
- FOCAL LENGTH CONSISTENCY — a wide-angle scene needs a wide-angle product perspective. A telephoto scene needs flattened product perspective.
- VANISHING POINT COHERENCE — the product's perspective lines must converge to the same vanishing points as the scene.
- SCALE COHERENCE — the product's apparent size relative to scene elements must be physically plausible.

7.3 Surface Interaction Simulation — NEW IN V3

Even with perfect lighting direction, a composited product looks pasted without these:

- AMBIENT OCCLUSION — darker shadows in crevices and at contact points where light is occluded
- ENVIRONMENTAL BOUNCE COLOR — the product picks up subtle color from nearby surfaces (a red flower next to the product adds a hint of red to the product's shadow side)
- CONTACT REFLECTION — if the product sits on a glossy surface, generate a faint reflection beneath it
- SUBSURFACE SCATTERING — for translucent materials (skincare gels, glass perfume bottles), light penetrates and scatters within

7.4 Material-Aware Relighting — NEW IN V3

Different materials react to light completely differently. A single relighting model produces wrong results for at least one material class.

Material Class	Lighting Behavior	Relight Approach
Plastic (matte)	Diffuse reflection, no specular	Soft directional shading, no highlights
Plastic (glossy)	Specular highlights, sharp edges	Match scene highlight direction precisely
Metallic	Highly reflective, mirrors environment	Synthesize environment reflection map
Glass / translucent	Light passes through, edge refraction	Render refraction edges, transmitted color
Fabric	Soft diffuse with subsurface scattering	Soft shading with slight glow on backlit edges
Ceramic	Matte with subtle specular	Diffuse + restrained specular highlight
Liquid	Refraction + surface reflection	Render meniscus, internal refraction, surface highlights

7.5 Composite vs Generate Decision Matrix

Product Type	Approach	Rationale
Cosmetics, packaged goods, electronics	Composite the real product	Fidelity is non-negotiable
Branded apparel	Composite where logo/print is critical	Logo accuracy required
Generic fashion / lifestyle	Native generation acceptable	Vibe matters more than exact match
Premium / hero campaigns	Two-pass hybrid	Best of both
Quick promotional posts	Native generation	Speed over perfection

8. Sensitive Content Protection (Hardened)

V2 used a -200 score penalty. V3 converts this to hard exclusion logic, adds ensemble moderation, and includes manual escalation rules.

8.1 The Hard Exclusion Rule

```
// V3: Hard exclusion, not weighted penalty function filterStylesForSensitiveContent(styles, context) { if
(!context.is_sensitive) { return styles; } // HARD EXCLUSION: remove any style with energy_level > 4 // or
minimal === false return styles.filter(style => style.energy_level <= 4 && style.minimal === true
&& !style.avoid_for_intents.includes('tribute') && !style.avoid_for_intents.includes('memorial') ); } // If no
styles remain after filtering, escalate to manual review // rather than fallback to a default style if
(filteredStyles.length === 0) { await escalateToManualReview(context); return { recommendation: null,
requires_human_review: true }; }
```

8.2 Ensemble Moderation

Single-classifier dependency is fragile. V3 uses ensemble moderation for sensitive content detection.

1. GPT-4o-mini classifies the prompt with high temperature variance (3 attempts)
2. If any of the 3 attempts returns `is_sensitive=true`, the content is treated as sensitive
3. Claude Sonnet performs a second-pass classification using different prompt phrasing
4. If GPT and Claude disagree, the content is treated as sensitive AND flagged for manual review
5. Specific trigger phrases (“passed away,” “in loving memory,” “RIP,” “death,” “funeral,” “mourning,” “condolences,” “accident,” “tragedy”) force `is_sensitive=true` regardless of classifier output

8.3 Uncertainty Thresholds

The classifier returns a confidence score. When confidence is below threshold, the system errs on the side of treating content as sensitive:

```
if (classification.confidence < 0.85 && classification.could_be_sensitive) { // Treat as sensitive when uncertain
context.is_sensitive = true; context.requires_review = true; }
```

8.4 Manual Escalation Rules

Cases that bypass automated style selection and require human review:

- Any post containing death-related trigger phrases
- Posts that mention disasters, accidents, or tragedies
- Religious mourning content (any faith)
- Posts about illness or hospitalization
- Any post the classifier flags with confidence below 0.85
- Any post where GPT and Claude disagree on sensitivity

Escalation flow: Jane responds to the user “This post needs special handling. I’ve flagged it for review by a human team member. We’ll be back within 30 minutes.” Meanwhile a Slack alert pings the URI Social team.

9. The Expanded Style Ontology

V2 used 3 dimensions for style metadata: intent, emotion, energy. Expert review correctly identified this as too shallow. V3 expands to 11 dimensions that better capture the actual dimensionality of visual branding.

9.1 The 11 Style Dimensions

Dimension	Spectrum	Example Tags
Energy	1-10	Tribute=1, Sale=10
Formality	luxury ↔ casual	Hotel campaign ↔ thrift vendor
Tone	playful ↔ serious	Meme ↔ corporate announcement
Density	maximalist ↔ minimalist	Layered editorial ↔ single-element studio
Treatment	documentary ↔ stylized	Behind-the-scenes ↔ fashion editorial
Subject Focus	human-centered ↔ product-centered	Founder portrait ↔ packshot
Visual Mode	cinematic ↔ graphic	Film-still ↔ poster design
Audience Register	premium ↔ meme-native	Luxury campaign ↔ viral meme
Composition	dense ↔ spacious	Information-heavy ↔ negative-space heavy
Emotion	specific from 10-word list	exciting, calm, urgent, intimate, bold, minimal, playful, serious, inspirational, nostalgic
Intent	specific from 12-action list	product_launch, sale, announcement, tribute, etc.

9.2 Updated Style Metadata Schema

```
{  slug: 'cinematic_founder',  name: 'Cinematic Founder Narrative',  // Aesthetic fields (used in prompt assembly)  quality_benchmark: 'Premium startup campaign visual. Cinematic documentary realism, shot on 85mm',  genre_description: 'Documentary realism with handheld documentary framing, soft halation, tungsten practicals',  composition_description: 'Documentary editorial framing with environmental wrapping and foreground occlusion',  lighting_description: 'Warm tungsten practical lighting. Motivated shadow direction. Soft window light bounce.',  texture: 'Kodak Portra tones. 35mm film grain. Soft highlight bloom.',  color_palette: 'Warm muted cinematic palette',  mood_words: ['driven', 'authentic'],  // V3: Expanded ontology (11 dimensions)  energy_level: 5,  formality: 'casual',  // luxury | casual  tone: 'serious',  // playful | serious  density: 'balanced',  // maximalist | minimalist | balanced  treatment: 'documentary',  // documentary | stylized  subject_focus: 'human',  // human-centered | product-centered  visual_mode: 'cinematic',  // cinematic | graphic  audience_register: 'premium',  // premium | meme-native  composition_density: 'balanced',  // dense | spacious | balanced  emotion: 'inspirational',  intent_tags: ['brand_story', 'storytelling'],  // Hard rules  minimal: false,  sensitive_content_safe: false,  // Industry alignment  industries: ['tech', 'professional_services', 'startup'], }
```

9.3 Why Higher Dimensionality Matters

With 11 dimensions, the recommendation engine can detect subtle mismatches V2 missed:

- A user asks for a “luxury” brand post but the recommended style is documentary-treatment. Mismatch detected on the formality dimension.
- A user wants a meme but the recommended style is premium audience-register. Caught on audience_register.
- A user wants a minimal announcement but the recommended style is maximalist. Caught on density.

The richer ontology lets the recommendation engine distinguish styles that previously looked similar on the limited 3-dimension scoring.

10. African Realism (Deepened)

V2 said “natural skin tones, visible individuality.” Expert review correctly noted this produces Western stock photography with dark skin rather than genuine African realism. V3 expands the vocabulary with specific local references.

10.1 The Updated African Realism Vocabulary

For African subjects, environments, and contexts, use these instead of generic descriptors:

Subjects

- Realistic Nigerian / Ghanaian / Kenyan / South African [subject role]
- Authentic African features and natural skin tone variation
- Visible individuality — not a generic stock face
- Realistic body proportions for African physique
- Natural hair textures (4A, 4B, 4C, locs, braids, twists, fades — specify)

Environments — Lagos / Nigerian Specific

- Authentic Lagos street scene with informal commercial signage
- Nigerian interior with practical Naija decor (woven mats, locally made furniture)
- Lagos commercial district with mixed formal and informal vendors
- Modern Lagos apartment interior — mid-tier middle-class realism
- Open-air Nigerian market with stall-based vendors
- Lagos roadside with danfo (yellow bus) elements
- Suya joint / mama-put atmosphere
- Owambe / traditional Nigerian event setting

Fashion Silhouettes

- Ankara prints with contemporary tailoring
- Modern Nigerian streetwear silhouettes
- Traditional Yoruba / Igbo / Hausa attire references where appropriate
- Aso ebi context for celebratory content
- Modest contemporary modest fashion
- Lagos office wear — sharp tailoring with Naija identity

Environmental Texture

- Tropical light quality — strong, often overhead, harsh midday or warm golden late afternoon
- Lagos atmospheric humidity in the light
- West African red earth color cast in environments
- Naija commercial signage typography and color palettes

10.2 What To Stop Doing

- Generic “African person” — too vague, defaults to stock photography
- “Traditional African clothing” — patronizing and inaccurate
- “African village” — reductive; Africa has cities, suburbs, modern interiors

- Using only generic Western stock photography lighting (clinical softbox)
- Assuming all African subjects are in rural or traditional contexts

10.3 The African Realism Block Template

Realistic [Nigerian / Lagos / Naija] [subject role: e.g., young entrepreneur, market vendor, fashion model], natural body language, authentic emotional realism, visible individuality, natural skin tone variation, [specific hair texture].
[Environment: e.g., authentic Lagos commercial street, modern Lagos apartment interior, open-air Nigerian market], lived-in realism with tropical light quality.

11. Brand Color Integration

Unchanged from V2 — the rules here were already correct.

BRAND: Subtly integrate the brand palette through environmental accents, wardrobe, props, reflections, gradients, or lighting cues. Primary palette: [descriptive color word] ([HEX]). Example: "deep magenta (#C2185B)" or "warm coral (#FF6B47)"

Always pair HEX with descriptive color word. HEX alone is weakly understood.

11.1 What To Avoid

- “Use brand color #C2185B throughout” — produces unnatural overlay
- “Apply brand filter” — produces flat tint
- “Dominant magenta” — overwhelms composition

11.2 What Works

- “Deep magenta accents in the environment”
- “Magenta reflections on the product”
- “Warm sunset palette with magenta undertones”

12. Text and Typography

Unchanged in principle from V2 — typography is composited in post-processing, not rendered by the image model. This decision was validated by expert review as one of the most commercially mature decisions in the system.

12.1 The Production Rule

HARD RULE

Image models cannot be trusted with typography. For production use at scale, typography MUST be composited in post-processing.

12.2 The Two-Pass Typography Pipeline

1. Image model generates the scene WITHOUT text. Prompt instructs: “Render NO text in the image.” Leave clean negative space at specified positions.
2. Server-side compositing renders text using the actual brand font file (.ttf or .otf) at exact size, weight, color, position.

12.3 What’s Missing (For Future Builds)

Expert review correctly noted that V2 thinks of typography as font rendering, when it should be:

- Layout intelligence — deciding WHERE text goes based on the image
- Hierarchy — headline vs subhead vs CTA visual weight
- Spacing rhythm — line height, letter spacing, paragraph spacing
- Typographic composition — text as part of overall composition

Current state: we composite text in pre-defined zones. Future state: we composite text with automated art direction. This is in the roadmap.

13. Composition Rules

13.1 Default Composition Patterns by Content Type

Content Type	Default Composition	Key Vocabulary
Product hero shot	Centered with environmental wrapping	foreground occlusion, atmospheric layering
Lifestyle / brand story	Documentary editorial framing	handheld documentary, natural asymmetry
Sale / promo	High-energy with clear text zones	graphic composition, bold contrast
Tribute / serious	Minimal, significant negative space	spacious composition, restrained
Behind-the-scenes	Candid documentary	handheld realism, off-camera flash
Food / restaurant	Top-down flatlay OR 45-degree editorial	macro texture realism, controlled diffusion
Fashion	Editorial portrait or full-body	runway editorial, off-camera flash aesthetic

13.2 Dynamic Motion by Product Category

Product Category	Motion Element
Cosmetics / skincare	Floating petals, soft mist, suspended liquid droplets
Food / drinks	Steam, splash, ingredient toss, falling garnish
Fashion	Wind through fabric, flowing motion, dynamic pose
Tech / hardware	Light streaks, particle effects, motion blur on accents
Jewelry	Sparkle, light refraction, gentle rotation
Hair / beauty services	Hair in motion, fabric sweep

14. The Visual Style Library

URI Social maintains 30+ visual styles. With the V3 expanded ontology, each style now stores 11 dimensions plus aesthetic fields.

Family	Count	Use Case
Editorial / Premium	8 styles	Product launches, brand campaigns, luxury
Lifestyle / Documentary	6 styles	Brand stories, behind-the-scenes, founder content
High-Energy / Sale	5 styles	Sales, promos, urgent campaigns
Atmospheric / Mood	5 styles	Storytelling, weekend content
Minimal / Respectful	4 styles	Tributes, serious announcements, professional
Art-Piece Posters	15 styles	Posters with logo + tagline integrated
Mobile Wallpapers	15 styles	9:16 background-only wallpapers
Meme Formats	12 styles	Engagement content, no brand logo
Video Styles	12 styles	Storyboard templates for video generation

15. Smart Style Recommendation (Updated for V3)

15.1 The Updated Flow

1. User types prompt: “promote our new lipstick”
2. Ensemble classification: GPT-4o-mini (×3 attempts) + Claude Sonnet
3. Sensitive content check: if positive, hard-exclude all non-minimal styles
4. Score every remaining style across 11 dimensions
5. Present three options: Recommended, Alternative (different vibe), User’s Default
6. User picks or taps “Let Jane decide”

15.2 The 11-Dimension Scoring

Each dimension contributes to the score. The total reflects how well the style matches the request across all dimensions, not just intent and emotion.

```
function scoreStyle(style, context, userPrefs, brand) {  let score = 0;      // Hard exclusions first  if
(context.is_sensitive && style.energy_level > 4) {    return { score: -Infinity, blocked: true };  }  if
(context.is_sensitive && !style.minimal) {            return { score: -Infinity, blocked: true };  }  // V3: 11-dimension
scoring  score += scoreEnergyMatch(style, context) * 1.5;  score += scoreFormalityMatch(style, context, brand) * 1.2;
score += scoreToneMatch(style, context) * 1.0;  score += scoreDensityMatch(style, context) * 0.8;  score +=
scoreTreatmentMatch(style, context) * 0.8;  score += scoreSubjectFocusMatch(style, context) * 1.0;  score +=
scoreVisualModeMatch(style, context) * 0.8;  score += scoreAudienceRegisterMatch(style, brand) * 1.2;  score +=
scoreCompositionDensityMatch(style, context) * 0.6;  score += scoreEmotionMatch(style, context) * 1.0;  score +=
scoreIntentMatch(style, context) * 1.5;  // User preference weighting  const userWeight = userPrefs[style.slug] || 1.0;
score *= userWeight;  // Brand industry alignment  if (style.industries.includes(brand.industry)) {    score +=
15;  }  return { score, blocked: false }; }
```

16. Image Quality Standards

Every generated image should exhibit these qualities:

- Depth separation — achieved via vocabulary (“cinematic depth separation,” “shallow focus,” “foreground occlusion”)

- Motivated lighting — light direction has a visible source in the scene
- Tonal richness — no pure black or pure white. Deep shadows and luminous highlights
- Material accuracy — via “photoreal material response,” “tactile detail,” “controlled specular”
- Atmospheric depth — film grain, haze, bokeh
- Composition coherence — the image reads as a single unified scene
- Emotional anchor — the image conveys a feeling matched to context
- Geometric correctness — perspectives align, proportions are physically plausible
- African subject authenticity — not Western stock with dark skin

17. What We Don’t Do

- Never generate competitor brand logos or products
- Never depict identifiable real people without permission
- Never generate sexual, gore, or graphic violence content
- Never generate content that could be confused with breaking news
- Never include the actual brand name as text in the image (text comes from the typography layer)
- Never use “leverage,” “elevate,” “seamless,” “game-changer” in text overlays
- Never apologize or explain limitations — the pipeline retries with a different prompt
- Never use the diluted vocabulary in Section 5.1

18. What Comes After V3

V3 brings the rulebook to a strong per-image quality standard. The next evolution is no longer about prompts — it’s about systems around the model. Items being prioritized for the next 6 months:

18.1 Campaign Consistency Memory (HIGHEST PRIORITY)

Store visual embeddings of every image generated for a brand. Compute a “brand visual signature” from the brand’s last N images. Use the signature as a soft constraint when generating new images. This is what makes Jane build visual identities, not just isolated images.

18.2 Automated Realism Scoring

After every generation, automatically score the output on: anatomy correctness, text cleanliness, product similarity, aesthetic consistency, lighting coherence. Use these scores to decide whether to deliver the image or regenerate.

18.3 Multi-Agent Critique Loop

Before delivering an image, run it through 4 critic agents: visual critic (composition, lighting, aesthetic), typography critic (text placement, hierarchy, readability), realism critic (anatomy, physics, material correctness), brand critic (consistency with brand signature, color, style).

18.4 Reference Image Conditioning

Allow users to upload mood boards that condition generation toward their preferred aesthetic. Combine text prompts with visual conditioning.

18.5 Cost and Latency Governance

Premium-tier generation (multi-pass, scored, refined) versus fast-tier generation (single-pass). Routing logic decides which tier based on the user's plan and the content type.

THE PROMPT CEILING

V3 brings the prompt layer to its practical limit. Further investment in prompt vocabulary will produce diminishing returns. The next leap in quality comes from systems AROUND the model: memory, retrieval, scoring, consistency. The prompt layer becomes one module of a larger visual generation platform.

Sharper vocabulary. Hard exclusions. Deeper compositing.

V3 is the floor. The systems we build next are the ceiling.